

Michael Mulugeta Demissie

Data Engineer | ML Infrastructure · ETL Pipelines · Cloud · Python

Washington, DC | michaelmulugetademissie@gmail.com | (202) 290-6215 | LinkedIn | GitHub | Portfolio

SUMMARY

Data Engineer with 3+ years of hands-on experience building production ETL pipelines, deploying ML models, and shipping full stack data applications. Skilled across the full data lifecycle, from ingestion and transformation to model serving and API delivery.

EDUCATION

The George Washington University | M.S. Data Analytics

Expected May 2027

Delhi Technological University | B.Tech Software Engineering

May 2025

TECHNICAL SKILLS

Data Engineering: Python, SQL, Airflow, Spark, Kafka, dbt, ETL/ELT, Data Modeling

Cloud & Infra: AWS (S3, Glue, Redshift, Lambda), Snowflake, BigQuery, Terraform, Docker

Databases: PostgreSQL, MySQL, Redis, Elasticsearch, MariaDB

ML & AI: PyTorch, TensorFlow, scikit-learn, HuggingFace, BERT, OpenAI API, MLflow

Full Stack: FastAPI, Flask, Node.js, Express, React, JavaScript, Bash, Git, CI/CD, JWT Auth

Visualization: Tableau, Power BI, Plotly, Dash, Grafana, Looker

EXPERIENCE

Data Engineer | Hasab AI

May 2025 – Apr 2026

Remote

- Architected Python ETL pipelines processing multilingual audio datasets across 10+ African and Middle Eastern languages, handling segmentation, noise filtering, and feature extraction.
- Built and deployed FastAPI services exposing structured NLP data to downstream speech and voice AI applications in underserved language markets.

Data Engineer | ParallelDots

Aug 2024 – Jun 2025

Delhi, India

- Engineered Python ETL pipelines ingesting retail shelf image metadata into MySQL, powering the ShelfWatch AI CV platform used by Fortune 500 FMCG clients globally.
- Built Pandas/NumPy preprocessing pipelines for CV model training and developed Flask REST APIs exposing processed data to production dashboards.

TECHNICAL PROJECTS

FinResolve: Complaints Data Pipeline & ML Routing System

Jan 2026 – May 2026

Python · Airflow · PostgreSQL · AWS S3 · Docker · FastAPI · DistilBERT · Render

- Built end-to-end ETL pipeline processing 1M+ CFPB records; engineered 812K training samples across 9 departments for ML routing.
- Fine-tuned DistilBERT for department classification and trained TF-IDF priority classifier achieving 88% accuracy; deployed via Docker on Render with JWT-auth FastAPI and role-based dashboards.

DC Metro Real-Time Arrivals Data Pipeline

May 2026

Node.js · PostgreSQL · Airflow · Docker · Express · Railway · WMATA API

- Built real-time ingestion pipeline polling WMATA API every 60s across all DC Metro stations, sustaining 580K+ rows/day with Airflow-scheduled rolling 24hr retention.
- Designed full architecture end-to-end: poller → PostgreSQL → Express API → live dashboard; containerized with Docker, deployed to production on Railway.